

# Proper way of handling threads in kernel?

 [stackoverflow.com/questions/10177641/proper-way-of-handling-threads-in-kernel](https://stackoverflow.com/questions/10177641/proper-way-of-handling-threads-in-kernel)

## Ask Question

I've seen bits of scattered information all around, but I can't seem to get to one final answer. How do you clean up a zombie thread in kernel?

Just to make sure, and produce a final correct way of handling threads in kernel, I would like to ask this question more broadly. **How do you create, terminate and clean up a thread in the Linux kernel?**

What I have so far is this:

```
thread_func:
    exited = 0;
    while (!must_exit)
        do stuff
    exited = 1;
    do_exit(0)

init_module:
    must_exit = 0;
    exited = 1;
    kthread_run(thread_func, ...)    /* creates and runs the thread */

cleanup_module:
    must_exit = 1;
    while (!exited)
        set_current_state(TASK_INTERRUPTIBLE);
        msleep(1);
    /* How do I cleanup? */
```

The closest thing I have found to the cleanup solution is `release_task`, but I didn't find anywhere talking about it. I imagined since the thread functions are `kthread_create`, `kthread_run` etc, there should be a `kthread_join` or `kthread_wait`, but there wasn't. `do_wait` also seemed likely, but it doesn't take a `struct task_struct *`.

Furthermore, I am not sure if `do_exit` is a good idea, or if at all necessary either. Can someone please come up with the minimum sketch of how a kthread should be created, terminated and cleaned up?

asked Apr 16 '12 at 16:03



Shahbaz

33.1k1080137

## 1 Answer 1

up vote 11 down vote accepted

One of the "right" ways to do this is to have your thread function check if it `kthread_should_stop` , and simply return if it does need to stop.

You don't need to call `do_exit` , and if you intend to `kthread_stop` it from the module exit function, you probably shouldn't.

You can see this by looking at the documentation for `kthread_create_on_node` in `kernel/kthread.c` (extract from Linux kernel 3.3.1):

```
/**
 * kthread_create_on_node - create a kthread.
 * @threadfn: the function to run until signal_pending(current).
 * @data: data ptr for @threadfn.
 * @node: memory node number.
 * @namefmt: printf-style name for the thread.
 *
 * Description: This helper function creates and names a kernel
 * thread. The thread will be stopped: use wake_up_process() to start
 * it. See also kthread_run().
 *
 * If thread is going to be bound on a particular cpu, give its node
 * in @node, to get NUMA affinity for kthread stack, or else give -1.
 * When woken, the thread will run @threadfn() with @data as its
 * argument. @threadfn() can either call do_exit() directly if it is a
 * standalone thread for which no one will call kthread_stop(), or
 * return when 'kthread_should_stop()' is true (which means
 * kthread_stop() has been called). The return value should be zero
 * or a negative error number; it will be passed to kthread_stop().
 *
 * Returns a task_struct or ERR_PTR(-ENOMEM).
 */
```

A "matching" comment is present for `kthread_stop` :

```
| If threadfn() may call do_exit() itself, the caller must ensure task_struct can't go away .
```

(And I'm not sure how you do that - probably holding on to the `struct_task` with a `get_task_struct` .)

If you walk the path of a thread creation you'll get something like:

```
kthread_create                                // macro in kthread.h
-> kthread_create_on_node                      // in kthead.c
-> adds your thread request to kthread_create_list
-> wakes up the kthreadd_task
```

`kthreadd_task` is set up in `init/main.c` in `reset_init` . It runs the `kthreadd` function (from `kthread.c` )

```
kthreadd // all in kthread.c
-> create_kthread
-> kernel_thread(kthread, your_kthread_create_info, ...)
```

And the `kthread` function itself does:

```
kthread
-> initialization stuff
-> schedule() // allows you to cancel the thread before it's actually started
-> if (!should_stop)
-> ret = your_thread_function()
-> do_exit(ret)
```

... So if `your_thread_function` simply returns, `do_exit` will be called with its return value. No need to do it yourself.

## Your Answer

---

Sign up or log in

---

Sign up using Google

Sign up using Facebook

Sign up using Email and Password

Post as a guest

---

By clicking "Post Your Answer", you acknowledge that you have read our updated [terms of service](#), [privacy policy](#) and [cookie policy](#), and that your continued use of the website is subject to these policies.

Not the answer you're looking for? Browse other questions tagged `c` `multithreading` `linux-kernel` `zombie-process` or ask your own question.

---

lang-c