

# textsearch infrastructure + skb\_find\_text()

 lwn.net/Articles/135159/

<b>From:</b>	Thomas Graf <tgraf@suug.ch>
<b>To:</b>	netdev@oss.sgi.com
<b>Subject:</b>	[RFC] textsearch infrastructure + skb_find_text()
<b>Date:</b>	Thu, 5 May 2005 01:40:36 +0200
<b>Cc:</b>	Pablo Neira <pablo@eurodev.net>
<b>Archive-link:</b>	<a href="#">Article</a> , <a href="#">Thread</a>

The patch below is a report on the current state of the textsearch infrastructure and its first user `skb_find_text()`. The textsearch is kept as simple as possible but advanced enough to handle non-linear data such as skb fragments. Unlike in many other approaches the text input is not seen as a single pointer but rather as a continuously called callback `get_text()` until 0 is returned allowing to search on any kind of data and to implement customized from-to limits.

The patch is separated into 3 parts, the first one being the textsearch infrastructure itself followed by a simple Knuth-Morris-Pratt implementation for reference. I'm also working on what could be called the smallest regular expression implementation ever but I left that out for now since it still has issues. Last but not least the function `skb_find_text()` written in a hurry and probably not yet correct but you should get the idea. From a userspace perspective the first user will be an ematch but writing it will be peanuts so I left it out for now.

Basically what it looks like right now is:

```
int pos;
struct ts_state;
struct ts_config *conf = textsearch_prepare("kmp", "hanky", 5, GFP_KERNEL, 1);

/* search for "hanky" at offset 20 until end of packet */
for (pos = skb_find_text(skb, 20, INT_MAX, conf, &state;
    pos >= 0;
    pos = textsearch_next(conf, &state)) {
    printk("Need a hanky? I found one at offset %d.\n", pos);
}

textsearch_put(conf);
kfree(conf);
```

You might wonder about the 1 given to `_prepare()`, it indicates whether to autoload modules because the ematches will need it to be able to drop `rtnl sem`.

The code is not tested and certainly not bug free yet but should compile.

Thoughts?

```
diff -X dontdiff -Nru linux-2.6.12-rc3.orig/include/linux/textsearch.h
linux-2.6.12-rc3/include/linux/textsearch.h
--- linux-2.6.12-rc3.orig/include/linux/textsearch.h 1970-01-01 01:00:00.000000000
+0100
+++ linux-2.6.12-rc3/include/linux/textsearch.h 2005-05-05 00:35:07.000000000
+0200
@@ -0,0 +1,169 @@
+#ifndef __LINUX_TEXTSEARCH_H
+#define __LINUX_TEXTSEARCH_H
+
+#ifdef __KERNEL__
+
+#include <linux/types.h>
+#include <linux/list.h>
+#include <linux/kernel.h>
+#include <linux/module.h>
+#include <linux/err.h>
+
+struct ts_config;
+
+/**
+ * struct ts_state - textsearch state
+ * @offset: current offset for next match
+ * @args: for persistent variables of get_text()
+ */
+struct ts_state
+{
+ int offset;
+ long args[6];
+};
+
+/**
+ * struct ts_ops - textsearch operations
+ * @name: name of search algorithm
+ * @init: called upon initialization to prepare a search
+ * @find: does the actual matching on a prepared configuration
+ * @owner: reference to algorithm module if existent
+ * @list: operations list we are on
+ */
+struct ts_ops
+{
+ const char *name;
+ struct ts_config * (*init)(const unsigned char *, size_t, int);
+ int (*find)(struct ts_config *,
+ struct ts_state *);
+ struct module *owner;
+ struct list_head list;
+};
+
+/**
+ * struct ts_config - textsearch configuration
+ * @ops: textsearch operations
+ * @get_text: callback to fetch text to search in
+ * @len: length of pattern
```

```

+ */
+struct ts_config
+{
+ struct ts_ops *ops;
+
+ /**
+  * get_text - return next chunk of text
+  * @offset: Number of bytes consumed by the matcher
+  * @dst: destination buffer
+  * @conf: search configuration
+  * @state: search state
+  *
+  * Gets called repeatedly until 0 is returned. Must assign a pointer
+  * to the start of the next chunk of text to *dst and return the
+  * length of the chunk or 0 if at the end. offset==0 indicates
+  * a new search. May store/read persistent values in state->args[].
+  */
+ int (*get_text)(int offset, unsigned char **dst,
+                  struct ts_config *conf,
+                  struct ts_state *state);
+
+ /**
+  * finish - called when the matching has been completed successful
+  *           or not.
+  * @conf: search configuration
+  * @state: search state
+  */
+ void (*finish)(struct ts_config *conf,
+                 struct ts_state *state);
+ int pattern_len;
+};
+
+/* Do not use this function directly */
+static inline int __textsearch_find(struct ts_config *conf,
+                                    struct ts_state *state)
+{
+ int ret = conf->ops->find(conf, state);
+
+ if (conf->finish)
+ conf->finish(conf, state);
+
+ return ret;
+}
+
+/**
+ * textsearch_find - search for a pattern in a text
+ * @conf: search configuration
+ * @state: search state
+ *
+ * Performs the actual search on the prepared configuration.
+ *
+ * Returns the position of first occurrence of the pattern or a
+ * negative number if no match was found.
+ */
+static inline int textsearch_find(struct ts_config *conf,
+                                   struct ts_state *state)

```

```

+{
+ state->offset = 0;
+ return __textsearch_find(conf, state);
+}
+
+/**
+ * textsearch_next - continue search for a pattern in a text
+ * @conf: search configuration
+ * @state: search state
+ *
+ * Continues a search looking for more occurrences of the pattern.
+ * You must call textsearch_find() to search the first occurrence
+ * in order to reset the state.
+ *
+ * Returns the position of next occurrence of the pattern or a
+ * negative number if no match was found.
+ */
+static inline int textsearch_next(struct ts_config *conf,
+    struct ts_state *state)
+{
+ state->offset += conf->pattern_len;
+ return __textsearch_find(conf, state);
+}
+
+/**
+ * textsearch_put - give back a textsearch configuration
+ * @conf: search configuration
+ *
+ * Releases all references of the configuration. Must be
+ * called prior to freeing the object.
+ */
+static inline void textsearch_put(struct ts_config *conf)
+{
+ if (conf->ops)
+ module_put(conf->ops->owner);
+}
+
+/**
+ * alloc_ts_config - allocate a textsearch configuration
+ * @payload: size of additional module specific data required
+ * @gfp_mask: allocation mask
+ *
+ * Returns a new, empty textsearch configuration or a ERR_PTR().
+ */
+static inline struct ts_config *alloc_ts_config(size_t payload, int gfp_mask)
+{
+ struct ts_config *conf;
+
+ conf = kmalloc(sizeof(*conf) + payload, gfp_mask);
+ if (conf == NULL)
+ return ERR_PTR(-ENOMEM);
+
+ memset(conf, 0, sizeof(*conf) + payload);
+ return conf;
+}
+
+extern int textsearch_register(struct ts_ops *);

```

```

+extern int textsearch_unregister(struct ts_ops *);
+extern struct ts_config *textsearch_prepare(const char *, const unsigned char
*,
+      size_t, int, int);
+extern int textsearch_find_continuous(struct ts_config *, struct ts_state *,
+      const unsigned char *, size_t);
+
+#endif /* __KERNEL__ */
+
+#endif
diff -X dontdiff -Nru linux-2.6.12-rc3.orig/lib/textsearch.c
linux-2.6.12-rc3/lib/textsearch.c
--- linux-2.6.12-rc3.orig/lib/textsearch.c 1970-01-01 01:00:00.000000000 +0100
+++ linux-2.6.12-rc3/lib/textsearch.c 2005-05-04 23:01:42.000000000 +0200
@@ -0,0 +1,241 @@
+/*
+ * lib/textsearch.c Generic text search interface
+ *
+ * This program is free software; you can redistribute it and/or
+ * modify it under the terms of the GNU General Public License
+ * as published by the Free Software Foundation; either version
+ * 2 of the License, or (at your option) any later version.
+ *
+ * Authors: Thomas Graf <tgraf@suug.ch>
+ *
+ * =====
+ *
+ * The textsearch infrastructure provides text searching facilities for both
+ * linear and non-linear data. Unlike many other
+ *
+ * Before a textsearch can be performed, a configuration must be created
+ * by calling textsearch_prepare() specifying the searching algorithm and
+ * the pattern to match. The returned configuration may then be used for
+ * an arbitrary amount of times and even in parallel as long as you
+ * provide a separate struct ts_state variable for every instance.
+ *
+ * The actual search is performed by either calling textsearch_find_text()
+ * for linear text or by providing an own get_text() implementation and
+ * calling textsearch_find(). Both functions will return the position
+ * of the first matching occurrence of the pattern. Subsequent matches
+ * may be retrieved via textsearch_next() irrespective of the linearity
+ * of the text.
+ *
+ * Once you're done using a configuration you must give back the
+ * allocated resources by calling textsearch_put().
+ *
+ * Example:
+ *   int pos;
+ *   struct ts_config *conf;
+ *   struct ts_state state;
+ *   const char *pattern = "chicken";
+ *   const char *example = "We dance the funky chicken";
+ *
+ *   conf = textsearch_prepare("kmp", pattern, strlen(pattern), GFP_KERNEL,
+ 1);
+ *   if (IS_ERR(conf)) {

```

```

+ *      err = PTR_ERR(conf);
+ *      goto errout;
+ *  }
+ *
+ *      pos = textsearch_find_text(conf, &state, example, strlen(example));
+ *      if (pos >= 0)
+ *          panic("Oh my god, dancing chickens at %d\n", pos);
+ *
+ *      textsearch_put(conf);
+ *      kfree(conf);
+ *
+ * =====
+ */
+
+#include <linux/config.h>
+#include <linux/module.h>
+#include <linux/types.h>
+#include <linux/string.h>
+#include <linux/err.h>
+#include <linux/textsearch.h>
+
+static LIST_HEAD(ts_ops);
+static DEFINE_RWLOCK(ts_mod_lock);
+
+static inline struct ts_ops *lookup_ts_algo(const char *name)
+{
+ struct ts_ops *o;
+
+ read_lock(&ts_mod_lock);
+ list_for_each_entry(o, &ts_ops, list) {
+ if (!strcmp(name, o->name)) {
+ if (!try_module_get(o->owner))
+ o = NULL;
+ read_unlock(&ts_mod_lock);
+ return o;
+ }
+ }
+ read_unlock(&ts_mod_lock);
+
+ return NULL;
+}
+
+/**
+ * textsearch_register - register a textsearch module
+ * @ops: operations lookup table
+ *
+ * This function must be called by textsearch modules to announce
+ * their presence. The given @ops must have %name set to a unique
+ * identifier and the callbacks find() and init() must be implemented.
+ *
+ * Returns -EEXISTS if another module has already registered with
+ * same name.
+ */
+int textsearch_register(struct ts_ops *ops)
+{
+ int err = -EEXIST;
+ struct ts_ops *o;

```

```

+
+ if (ops->name == NULL || ops->find == NULL || ops->init == NULL)
+   return -EINVAL;
+
+ write_lock(&ts_mod_lock);
+ list_for_each_entry(o, &ts_ops, list)
+   if (!strcmp(ops->name, o->name))
+     goto errout;
+
+ list_add_tail(&ops->list, &ts_ops);
+ err = 0;
+errout:
+ write_unlock(&ts_mod_lock);
+ return err;
+}
+
+/**
+ * textsearch_unregister - unregister a textsearch module
+ * @ops: operations lookup table
+ *
+ * This function must be called by textsearch modules to announce
+ * their disappearance for examples when the module gets unloaded.
+ * The @ops parameter must be the same as the one during the
+ * registration.
+ *
+ * Returns -ENOENT if no matching textsearch registration was found.
+ */
+int textsearch_unregister(struct ts_ops *ops)
+{
+   int err = 0;
+   struct ts_ops *o;
+
+   write_lock(&ts_mod_lock);
+   list_for_each_entry(o, &ts_ops, list) {
+     if (o == ops) {
+       list_del(&o->list);
+       goto out;
+     }
+   }
+
+   err = -ENOENT;
+out:
+   write_unlock(&ts_mod_lock);
+   return err;
+}
+
+static inline int get_linear_text(int offset, unsigned char **text,
+    struct ts_config *conf,
+    struct ts_state *state)
+{
+   unsigned char *string = (unsigned char *) state->args[0];
+   size_t len = state->args[1];
+
+   if (offset < len) {
+     *text = string + offset;
+     return len - offset;
+   }

```

```

+ } else
+ return 0;
+}
+
+/**
+ * textsearch_find_continuous - search a pattern in continuous/linear text
+ * @conf: search configuration
+ * @state: search state
+ * @text: text to search in
+ * @len: length of text
+ *
+ * A simplified version of textsearch_find() for continuous/linear text.
+ * Call textsearch_next() to retrieve subsequent matches.
+ *
+ * Returns the position of first occurrence of the pattern or a
+ *      negative number if no match was found.
+ */
+int textsearch_find_continuous(struct ts_config *conf, struct ts_state
+*state,
+    const unsigned char *text, size_t len)
+{
+    conf->get_text = get_linear_text;
+    state->args[0] = (long) text;
+    state->args[1] = len;
+
+    return textsearch_find(conf, state);
+}
+
+/**
+ * textsearch_prepare - Prepare a text search
+ * @algo: name of search algorithm
+ * @pattern: pattern data
+ * @len: length of pattern
+ * @gfp_mask: allocation mask
+ * @autoload: bool indicating whether to autoload modules
+ *
+ * Looks up the search algorithm module and creates a new textsearch
+ * configuration for the specified pattern. Upon completion all
+ * necessary refcnts are held and the configuration must be put back
+ * using textsearch_put() after usage.
+ *
+ * Note: The format of the pattern may not be compatible between
+ *      the various search algorithms.
+ *
+ * Returns a new textsearch configuration according to the specified
+ *      parameters or a ERR_PTR().
+ */
+struct ts_config *textsearch_prepare(const char *algo,
+    const unsigned char *pattern, size_t len,
+    int gfp_mask, int autoload)
+{
+    int err = -ENOENT;
+    struct ts_config *conf;
+    struct ts_ops *ops;
+
+    ops = lookup_ts_algo(algo);
+
+#ifdef CONFIG_KMOD

```



```

+ /* Why not always autoload you may ask. Some users may be
+  * in a situation where requesting a module may deadlock,
+  * especially when the module is located on a NFS mount. */
+ if (ops == NULL && autoload) {
+   request_module("ts_%s", algo);
+   ops = lookup_ts_algo(algo);
+ }
+ #endif
+
+ if (ops == NULL)
+   goto errout;
+
+ conf = ops->init(pattern, len, gfp_mask);
+ if (IS_ERR(conf)) {
+   err = PTR_ERR(conf);
+   goto errout;
+ }
+
+ conf->ops = ops;
+ return conf;
+
+errout:
+ if (ops)
+   module_put(ops->owner);
+
+ return ERR_PTR(err);
+}
+
+EXPORT_SYMBOL(textsearch_register);
+EXPORT_SYMBOL(textsearch_unregister);
+EXPORT_SYMBOL(textsearch_prepare);
+EXPORT_SYMBOL(textsearch_find_continuous);
diff -X dontdiff -Nru linux-2.6.12-rc3.orig/lib/ts_kmp.c
linux-2.6.12-rc3/lib/ts_kmp.c
--- linux-2.6.12-rc3.orig/lib/ts_kmp.c 1970-01-01 01:00:00.000000000 +0100
+++ linux-2.6.12-rc3/lib/ts_kmp.c 2005-05-05 00:32:28.000000000 +0200
@@ -0,0 +1,108 @@
+/*
+ * lib/ts_kmp.c  Knuth-Morris-Pratt text search implementation
+ *
+ * This program is free software; you can redistribute it and/or
+ * modify it under the terms of the GNU General Public License
+ * as published by the Free Software Foundation; either version
+ * 2 of the License, or (at your option) any later version.
+ *
+ * Authors: Thomas Graf <tgraf@suug.ch>
+ */
+
+ #include <linux/config.h>
+ #include <linux/module.h>
+ #include <linux/types.h>
+ #include <linux/string.h>
+ #include <linux/textsearch.h>
+
+ #define TS_KMP_PREFIX_TBL(conf) \
+ ((unsigned int *) ((unsigned char *) (conf) \

```

```

+     + sizeof(struct ts_config)))
+
+#define TS_KMP_PATTERN(conf) \
+ ((unsigned char *) TS_KMP_PREFIX_TBL(conf) \
+   + (conf->pattern_len * sizeof(unsigned int)))
+
+static int kmp_find(struct ts_config *conf, struct ts_state *state)
+{
+ int i, q = 0, consumed = state->offset;
+ unsigned char *text, *pattern = TS_KMP_PATTERN(conf);
+ unsigned int *prefix_tbl = TS_KMP_PREFIX_TBL(conf);
+ size_t text_len, pattern_len = conf->pattern_len;
+
+ for (;;) {
+   text_len = conf->get_text(consumed, &text, conf, state);
+
+   if (text_len == 0)
+     break;
+
+   for (i = 0; i < text_len; i++) {
+     while (q > 0 && pattern[q] != text[i])
+       q = prefix_tbl[q - 1];
+     if (pattern[q] == text[i])
+       q++;
+     if (q == pattern_len) {
+       state->offset = consumed + i - pattern_len + 1;
+       return state->offset;
+     }
+   }
+
+   consumed += text_len;
+ }
+
+ return -1;
+
+}
+
+static inline void compute_prefix_tbl(const unsigned char *pattern, size_t
len,
+   unsigned int *prefix_tbl)
+{
+ unsigned int k, q;
+
+ for (k = 0, q = 1; q < len; q++) {
+   while (k > 0 && pattern[k] != pattern[q])
+     k = prefix_tbl[k-1];
+   if (pattern[k] == pattern[q])
+     k++;
+   prefix_tbl[q] = k;
+ }
+}
+
+static struct ts_config *kmp_init(const unsigned char *pattern, size_t len,
+   int gfp_mask)
+{
+ struct ts_config *conf;
+

```

```

+ conf = alloc_ts_config(len + (len * sizeof(int)), gfp_mask);
+ if (IS_ERR(conf))
+ return conf;
+
+ conf->pattern_len = len;
+ memcpy(TS_KMP_PATTERN(conf), pattern, len);
+ compute_prefix_tbl(pattern, len, TS_KMP_PREFIX_TBL(conf));
+
+ return conf;
+}
+
+static struct ts_ops kmp_ops = {
+ .name = "kmp",
+ .find = kmp_find,
+ .init = kmp_init,
+ .owner = THIS_MODULE,
+ .list = LIST_HEAD_INIT(kmp_ops.list)
+};
+
+static int __init init_kmp(void)
+{
+ return textsearch_register(&kmp_ops);
+}
+
+static void __exit exit_kmp(void)
+{
+ textsearch_unregister(&kmp_ops);
+}
+
+MODULE_LICENSE("GPL");
+
+module_init(init_kmp);
+module_exit(exit_kmp);
diff -X dontdiff -Nru linux-2.6.12-rc3.orig/net/core/skbuff.c
linux-2.6.12-rc3/net/core/skbuff.c
--- linux-2.6.12-rc3.orig/net/core/skbuff.c 2005-04-30 20:22:24.000000000
+0200
+++ linux-2.6.12-rc3/net/core/skbuff.c 2005-05-05 00:35:50.000000000 +0200
@@ -1502,6 +1502,80 @@
     skb_split_no_header(skb, skb1, len, pos);
 }

+static int get_skb_text(int offset, unsigned char **text,
+ struct ts_config *conf, struct ts_state *state)
+{
+ /* args[0]: lower limit
+ * args[1]: upper limit
+ * args[2]: skb
+ * args[3]: fragment index
+ * args[4]: current fragment data buffer
+ * args[5]: octets consumed up to previous fragment */
+ int from = state->args[0], to = state->args[1];
+ struct sk_buff *skb = (struct sk_buff *) state->args[2];
+ int limit = min_t(int, skb_headlen(skb), to);
+ int real_offset = offset + from;
+ skb_frag_t *f;

```

```

+
+ if (!skb_is_nonlinear(skb)) {
+   if (real_offset < limit) {
+linear:
+   *text = skb->data + real_offset;
+   return limit - real_offset;
+ }
+
+ return 0;
+ }
+
+ if (real_offset < limit)
+   goto linear;
+
+next_fragment:
+ f = &skb_shinfo(skb)->frags[state->args[3]];
+ limit = min_t(int, f->size + state->args[5], to);
+
+ if (!state->args[4])
+   state->args[4] = (long) kmap_skb_frag(f);
+
+ if (real_offset < limit) {
+   *text = (unsigned char *) state->args[4] + f->page_offset +
+   (real_offset - (int) state->args[5]);
+   return limit - real_offset;
+ }
+
+ kunmap_skb_frag((void *) state->args[4]);
+ state->args[3]++;
+ state->args[4] = (long) NULL;
+ state->args[5] += f->size;
+
+ if (state->args[3] >= skb_shinfo(skb)->nr_frags)
+   return 0;
+
+ goto next_fragment;
+}
+
+static void get_skb_text_finish(struct ts_config *conf, struct ts_state
+*state)
+{
+ if (state->args[4])
+   kunmap_skb_frag((void *) state->args[4]);
+}
+
+int skb_find_text(struct sk_buff *skb, int from, int to,
+ struct ts_config *config, struct ts_state *state)
+{
+ config->get_text = get_skb_text;
+ config->finish = get_skb_text_finish;
+ state->args[0] = from;
+ state->args[1] = to;
+ state->args[2] = (long) skb;
+ state->args[3] = 0;
+ state->args[4] = 0;
+ state->args[5] = skb_headlen(skb);
+
+

```

```

+ return textsearch_find(config, state);
+}
+
+
void __init skb_init(void)
{
    skbuff_head_cache = kmem_cache_create("skbuff_head_cache",
@@ -1540,3 +1614,4 @@
    EXPORT_SYMBOL(skb_unlink);
    EXPORT_SYMBOL(skb_append);
    EXPORT_SYMBOL(skb_split);
+EXPORT_SYMBOL(skb_find_text);
diff -X dontdiff -Nru linux-2.6.12-rc3.orig/include/linux/skbuff.h
linux-2.6.12-rc3/include/linux/skbuff.h
--- linux-2.6.12-rc3.orig/include/linux/skbuff.h 2005-04-30 20:22:19.000000000
+0200
+++ linux-2.6.12-rc3/include/linux/skbuff.h 2005-05-05 00:12:46.000000000
+0200
@@ -28,6 +28,7 @@
#include <linux/poll.h>
#include <linux/net.h>
#include <net/checksum.h>
+#include <linux/textsearch.h>

#define HAVE_ALLOC_SKB /* For the drivers to know */
#define HAVE_ALIGNABLE_SKB /* Ditto 8) */
@@ -1186,6 +1187,9 @@
extern void      skb_split(struct sk_buff *skb,
        struct sk_buff *skb1, const u32 len);

+extern int skb_find_text(struct sk_buff *skb, int from, int to,
+    struct ts_config *config, struct ts_state *state);
+
+static inline void *skb_header_pointer(const struct sk_buff *skb, int offset,
        int len, void *buffer)
{

```